

000 ARITHMETIC-BENCH: EVALUATING MULTI-STEP 001 REASONING IN LLMs WITH BASIC ARITHMETIC 002

003 **Anonymous authors**
004

005 Paper under double-blind review
006

007 ABSTRACT 008

009 We propose Arithmetic-Bench, a benchmark designed to evaluate the multi-step
010 reasoning ability of large language models (LLMs) through basic arithmetic
011 operations. The benchmark covers fundamental mathematical operations such as
012 addition, subtraction, multiplication, and division, while also incorporating sub-
013 tasks like copying, reversing, counting, and base conversion. Experimental results
014 show that the accuracy of current LLMs drops sharply when performing arithmetic
015 operations involving more than 10 digits, implying a failure of generalization in
016 multi-step reasoning. We further analyze the root causes of these failures. While
017 LLMs can achieve a certain degree of arithmetic generalization through training
018 on limited-length sequences, they fail to generalize to arbitrary lengths. This is due
019 to the inherent complexity of arithmetic tasks: achieving true arithmetic general-
020 ization cannot rely on memorization alone but requires the acquisition of genuine
021 reasoning mechanisms. Compared to other math benchmarks, Arithmetic-Bench
022 provides a simple and fair framework. Because the tasks are purely synthetic, they
023 are easy to generate and largely free from human biases. We believe that arith-
024 metic tasks are both fundamental and necessary for advancing reasoning models,
025 and Arithmetic-Bench offers a principled way to evaluate them.
026

027 1 INTRODUCTION 028

029 The rapid development of large language models (LLMs) has led to significant progress in natural
030 language understanding and generation. However, despite their strong performance on existing
031 reasoning benchmarks such as AIME Veeraboina (2023), GSM8K Cobbe et al. (2021), and MATH
032 Hendrycks et al. (2021), these models often struggle with basic arithmetic tasks. This inconsistency
033 raises critical questions about the nature of reasoning in LLMs: do they truly possess multi-step
034 reasoning capabilities, or are they merely performing pattern matching based on training data?
035

036 1.1 DISADVANTAGES OF MATH BENCHMARKS 037

038 There are a lot of existing math-related datasets and benchmarks. However, in practical applications
039 involving reasoning models, we have observed several limitations in these math benchmarks.
040

041 **Hard to Collect.** Although a large number of math problems are available online, their difficulty
042 levels and coverage are difficult to control precisely, often resulting in datasets with uneven qual-
043 ity and distribution. Manually creating problems is costly and labor-intensive, while the reliability
044 of model-generated problems remains uncertain. In addition, manually collected problems are in-
045 evitably subject to human biases.

046 **Hard to Decontaminate (Easy to Cheat).** Given the vast amount of data on the internet, it is
047 almost inevitable that identical or highly similar problems already exist. Furthermore, it is difficult
048 to prevent individuals or organizations from intentionally training models on benchmark data to
049 inflate performance.

050 **Hard to Evaluate.** Evaluation poses significant challenges. For open-ended problems, such as
051 proofs, relying on models for evaluation is unreliable and vulnerable to hacking. For problems with
052 definitive answers, such as computational tasks, formatting problems necessitate complex pattern-
053 matching methods to verify correctness, which are error-prone and cause fluctuations in evaluation
results. This makes it difficult to determine whether a model's reasoning ability has truly improved.

054 **Hard to Scale.** Constructing a smooth difficulty progression is highly challenging. Some problems
 055 are difficult due to reliance on obscure knowledge, while others require multi-step reasoning. Con-
 056 sequently, certain problems primarily test memory rather than reasoning. Since these two types of
 057 difficulty differ in nature, they cannot be directly compared or quantified.
 058

059 1.2 ADVANTAGES OF ARITHMETIC BENCHMARKS
 060

061 In contrast to complex mathematical operations, arithmetic operations serve as a natural testbed for
 062 reasoning because they are deterministic, require structured multi-step execution, and have clear
 063 correctness criteria.

064 **Easy to Collect.** Generating arithmetic expressions is very straightforward, and results obtained
 065 through a calculator are guaranteed to be correct, ensuring the quality of the problems. More impor-
 066 tantly, these problems are purely synthetic, which greatly reduces human bias.
 067

068 **Easy to Decontaminate (Hard to Cheat).** Thanks to the nature of large numbers, no special filter-
 069 ing is required. Memorizing answers provides no advantage in large-number arithmetic: even if one
 070 memorizes all two-digit multiplications, they would cover only 1% of three-digit multiplications.
 071 Moreover, brute-force memorization inevitably leads to forgetting.

072 **Easy to Evaluate.** There is no ambiguity: evaluation can be performed directly with a simple
 073 check (e.g., $a \in b$), without additional prompts, since all mainstream models already know basic
 074 arithmetic and the numbers do not suffer from formatting issues.

075 **Easy to Scale.** Arithmetic tasks can be scaled to arbitrary digit lengths and varying levels of com-
 076 plexity, which creates a continuous difficulty curve. This makes it possible to evaluate a model’s
 077 true reasoning ability, going beyond mere memorization.
 078

079 1.3 MEMORIZATION VS GENERALIZATION
 080

081 We further raise the following two key questions and address them using Arithmetic-Bench.

082 **How can we verify whether the improvements of existing LLMs on reasoning benchmarks
 083 may come from memorizing the answers?**

085 It is possible to train a model on a finite math benchmark dataset and then achieve very high scores
 086 on that benchmark. However, even if we train on a multiplication benchmark dataset, the model still
 087 cannot achieve high scores on the multiplication benchmark, because the multiplication benchmark
 088 is randomly generated from a space that far exceeds the model’s capacity limit.

089 **How can we construct tasks that cannot be solved by memorizing the answers?**

090 The information required to fully memorize long multiplication numbers is infinite, whereas the
 091 information needed to memorize the rules of multiplication is finite. The space of multiplication is
 092 so large that, unlike Olympiad math problems which are finite, it cannot be completely memorized;
 093 only by understanding the rules of multiplication can true generalization be achieved.
 094

095 1.4 PROXY METRIC
 096

097 In the field of image generation, text rendering seems like a minor skill, but the Nano Banana Google
 098 (2025) team treated it as an important metric. Text is highly structured, and even small stroke errors
 099 are obvious, making it a strict test of precision. Mastering text forces the model to control structure
 100 and detail at the pixel level, which then improves general quality. By using text rendering as a proxy
 101 metric, the team showed how optimizing for a highly demanding, low-tolerance subtask can push
 102 models to develop transferable skills that enhance broader performance.

103 Arithmetic-Bench is also this type of task, requiring models to have stable and precise reasoning
 104 abilities, which are necessary for solving truly complex problems, such as Fermat’s Last Theorem.
 105 The formal proof of Fermat’s Last Theorem contains tens of thousands of lines of Lean code Buzzard
 106 & contributors (2025); de Moura & Ullrich (2021), and even without considering the details of each
 107 step, it still represents an extremely long chain of reasoning. Therefore, we believe that Arithmetic-
 Bench is suitable as a proxy metric for mathematical reasoning.

108 Whether a model truly has reasoning ability is a rather vague question, but if it can handle large-
 109 number arithmetic well, it can then be considered to possess a certain level of reasoning ability.
 110

111 The contributions of this paper are as follows:

- 112
- 113 • We introduce **Arithmetic-Bench**, a benchmark consisting of basic arithmetic tasks de-
 114 signed to evaluate models’ multi-step reasoning and computational skills.
 - 115 • We provide a theoretical analysis of the connection between arithmetic tasks and reasoning
 116 ability, and empirically demonstrate their correlation.
 - 117 • We benchmark multiple mainstream models, showing that current models perform poorly
 118 on these tasks, underscoring the need to improve multi-step reasoning capabilities.
- 119

120 2 RELATED WORK

121

122 2.1 MATH BENCHMARK

123

124 AIME Veeraboina (2023) (American Invitational Mathematics Examination) consists of
 125 competition-level math problems, covering advanced algebra, number theory, combinatorics, and
 126 geometry. It includes 15 problems per year before 2000 and 30 problems per year thereafter.

127 MATH Hendrycks et al. (2021) contains high-school level math problems spanning algebra, cal-
 128 culus, number theory, and more, totaling 12,500 problems. Large language models (LLMs) still
 129 struggle on some of these problems, particularly those requiring multi-step reasoning, achieving
 130 only moderate accuracy.

131 CMATH Wei et al. (2023) is a dataset of Chinese elementary school math word problems, compris-
 132 ing 1.7k problems with detailed annotations sourced from real workbooks and exams.

134 GSM8K Cobbe et al. (2021) (Grade School Math 8K) is a set of 8,000 elementary-level math prob-
 135 lems. Current LLMs can perform well on this benchmark, achieving over 97% accuracy through
 136 prompt engineering Zhong et al. (2024).

137 GSM-Symbolic Mirzadeh et al. (2024) is a variant of GSM8K in which numbers are replaced with
 138 random values. The resulting performance drop indicates that models may rely on memorized num-
 139 bers and patterns.

140 These benchmarks cover a broad spectrum from elementary arithmetic to advanced competi-
 141 tion-level mathematics.

143 2.2 ARITHMETIC BENCHMARK

144

145 Benchmarks focusing specifically on arithmetic, such as Math401 Yuan et al. (2023) and the arith-
 146 metic subset of BIG-Bench Srivastava et al. (2023), evaluate basic operations but have two key
 147 limitations: (i) potential memorization due to fixed datasets, and (ii) limited length generalization,
 148 since most problems involve numbers with fewer than ten digits. To address these issues, some ap-
 149 proaches use synthetic math games Kurtic et al. (2024), though these often require complex rules and
 150 careful prompt design. In contrast, Arithmetic-Bench provides a simpler framework for evaluating
 151 arithmetic reasoning with controlled difficulty and sequence length.

152 2.3 ARITHMETIC REASONING BASED ON DEEP LEARNING

153

154 Early works, including Neural GPU Łukasz Kaiser & Sutskever (2016) and Neural Turing Machine
 155 Graves et al. (2014), improved algorithm execution by designing specialized architectures such as
 156 recursive convolutional networks and memory modules. More recent methods, such as Goat Liu
 157 & Low (2023) and MathGLM Yang et al. (2023), train LLMs on carefully constructed arithmetic
 158 datasets, while other approaches, like Scratchpad Nye et al. (2021), leverage techniques such as
 159 chain-of-thought (CoT) reasoning Wei et al. (2022) and curriculum learning Bengio et al. (2009).
 160 These methods enhance arithmetic performance for numbers within certain digit lengths, but they
 161 generally fail to generalize to longer sequences, highlighting the challenge of length extrapolation
 in arithmetic reasoning.

162

3 ARITHMETIC-BENCH

163

164

3.1 CAPACITY

165

166 Mathematical reasoning is based on axioms and deductive rules. Here, axioms provide fundamental
167 assumptions, and deductive rules specify how to derive new facts from existing ones. A proposition
168 is considered correct if it can be derived from the axioms. The key difference between reasoning and
169 common sense lies in the number of inference steps: common sense typically requires only a single
170 step, whereas reasoning involves multiple iterative steps. Following Zhou (2025), reasoning models
171 are characterized by producing intermediate reasoning tokens before generating the final output.

172 From these observations, we propose the following definitions:

174 **Definition 1.** Reasoning is the iterative application of operations on finite information, where each
175 operation transforms known information into new information.

177 **Definition 2.** Arithmetic is a special case of reasoning, where the operations are derived from a
178 finite lookup table of number operations.

179 Clearly, arithmetic over natural numbers satisfies Definition 1 and is therefore a form of reasoning.

180 A task with finite information can be fully learned by memorizing all cases, provided the model has
181 sufficient capacity. Here, capacity refers to the total amount of information that a model can store
182 or represent in its parameters. More concretely, if a model has N parameters and each parameter
183 can store approximately c bits of information independently, then the model capacity is roughly
184 $C = N \cdot c$ bits. When the information content of a task exceeds the model's capacity, the task
185 becomes unlearnable due to inevitable forgetting. This is formalized by the following principle:

187 **Theorem 1.** A container with capacity a cannot hold information exceeding a .

189 For example, a model with 400 parameters can store the 9×9 multiplication table; a model with
190 20,000 parameters can fully memorize the first 10,000 digits of π . In contrast, a model with 10,000
191 parameters can memorize only about 70% of these digits, regardless of training duration. This is
192 coincidence with the fact that current language models can and only can store 2 bits of knowledge
193 per parameter Allen-Zhu & Li (2024).

194 Next, we relate arithmetic performance to general reasoning ability.

196 **Theorem 2.** If a model cannot learn an arithmetic problem, it cannot learn a reasoning problem of
197 equivalent complexity.

198 **Proof.** Any reasoning task can be encoded as an equivalent arithmetic problem by mapping basic
199 operations to numbers. If a model can solve this arithmetic problem, it can solve the corresponding
200 reasoning task. By Theorem 1, if the model cannot solve the arithmetic problem, it lacks sufficient
201 capacity to represent the necessary information, and therefore cannot learn any reasoning problem
202 of equal or greater complexity.

203 Computational stability can be analyzed similarly. Suppose each operation introduces a small error
204 ϵ , and the task can tolerate an expected error δ . Then, only a limited number of operations can be
205 performed before the accumulated error exceeds δ , defining the model's computational capacity.

206 A model can reliably complete a reasoning task only if both its information capacity and compu-
207 tational capacity are sufficient. Notably, increasing the number of digits in arithmetic primarily
208 challenges computational capacity rather than information capacity. Therefore, benchmarks like
209 Arithmetic-Bench can probe reasoning ability beyond the limits of information storage by evalua-
210 ting tasks that require extensive computation.

212

3.2 ERROR ACCUMULATION

213

214 Assuming the model has sufficient information capacity and fully understands the reasoning rules,
215 errors may still occur due to probabilistic predictions. To mitigate accumulated errors during itera-
216 tive reasoning, verification strategies can be employed. Let the probability of making a mistake in a

single computation be p , assuming independent computations. Without verification, the probability of submitting an incorrect result is

$$P_{\text{error, no verification}} = p.$$

If one additional independent verification is performed, an undetected error occurs only if the first computation is wrong, the verification is also wrong, and the two errors coincide exactly. Let q denote the conditional probability that two independent errors yield the same incorrect result ($0 < q < 1$). Then, the probability of an undetected error under verification is

$$P_{\text{error, verification}} = p^2 q.$$

Since $0 < p < 1$ and $0 < q < 1$, it follows that

$$P_{\text{error, verification}} = p^2 q < p = P_{\text{error, no verification}}.$$

Therefore, verification reduces the probability of undetected errors.

For instance, if $p = 0.1$ and $q = 0.1$, the error probability without verification is 10%, while with verification it decreases to 1.0%, representing a reduction by two orders of magnitude. This illustrates that, for reasoning tasks, implementing verification can be more effective than merely increasing the number of reasoning steps or output tokens.

3.3 DESIGN

Arithmetic-Bench is a dynamic benchmark that generates arithmetic problems of varying lengths and complexities. It includes binary and unary operations as shown in Table 1 and Table 2:

Each problem is randomly generated to ensure that tasks cannot be memorized, and all require multi-step iterative operations. The benchmark only requires the model to have basic mathematical knowledge, without relying on any advanced theorems to eliminate the influence of memorized knowledge on problem difficulty. Prompts are kept as simple as possible to minimize the influence of prompt- or instruction-following abilities on the results. Evaluation is performed directly using ‘a in b’, which is simple and accurate. Different models may format their outputs differently; for example, DeepSeek outputs answers in ‘\boxed{}’ and GPT prefers bold answers using ‘**’, but ‘a in b’ can match any similar format. If a model produces the correct intermediate result during the process, it indicates that the model’s reasoning can reach the final answer. Since the probability of guessing large-number results correctly in the middle of the process is extremely low, this does not compromise fairness. Some models, such as DeepSeek, may output answers separated by symbols like ‘’, ‘/!’. We remove all such symbols from the results before matching the answers. The steps for decimal addition and multiplication are basically the same as for integers, differing only in the decimal point shift. Therefore, only integer operations are considered. To ensure comparable computational complexity between division and multiplication and to avoid results that are too small, we perform division of $2n$ -digit numbers by n -digit numbers. Modular operations (mod) and division steps are essentially the same, so only division is considered. Exponentiation (pow) is too computationally expensive and is therefore not considered. We primarily evaluate the model’s arithmetic performance from the following two dimensions:

Accuracy

Full-match accuracy is used, without considering digit-wise accuracy, since in large-number operations models often produce outputs with incorrect digit lengths, making alignment with the correct answer difficult.

Length Generalization Curve

This curve illustrates the relationship between model accuracy and the number of digits in the input. It provides insight into how well the model can generalize to longer sequences, indicating its computational capacity.

3.4 PROMPT

The prompts for different tasks are shown in Table 3. They are designed to be as simple as possible to enhance readability for both humans and models.

Table 1: Main tasks: standard arithmetic tasks.

Task	Description
Add, Sub, Mul, Div	Integer addition, subtraction, multiplication, and division of 2 numbers, where both operands are n -digit integers. These tasks evaluate the model’s ability to perform standard arithmetic operations and its accuracy and consistency across multiple digits.
Add_1, Sub_1, Mul_1, Div_1	Integer addition, subtraction, multiplication, and division of 2 numbers, where one operand is an n -digit integer and the other is a single-digit integer. Since $n \times n$ multiplication can be decomposed into multiple $n \times 1$ multiplications, this task is used for evaluation.

Table 2: Sub tasks: basic operations related to arithmetic.

Task	Description
Copy	In multi-step reasoning, the model often needs to repeat operands multiple times. This task evaluates the model’s ability to correctly copy values within a reasoning chain.
Rev, Space	Data representation significantly affects performance Lee et al. (2023). For instance, columnar (vertical) arithmetic is written from right to left, which can hinder next-token prediction. Reversal and splitting operations, such as little-endian storage or separating numbers into individual characters, are evaluated to test the model’s adaptability to different representations.
Count, Len	Models sometimes produce outputs of incorrect length. These tasks test the model’s counting ability. Since the answers are small numbers, parentheses are used to prevent models from accidentally guessing the correct answer during generation.
Box	Some operations, like count and len, require formatted output. This task evaluates the model’s ability to correctly insert parentheses as a formatting symbols.
B2D, D2B	Neural GPU Łukasz Kaiser & Sutskever (2016) has shown better performance in binary than decimal. These tasks evaluate the model’s ability to convert between binary and decimal representations.

3.5 GENERATION

The generation of Arithmetic-Bench is very straightforward: two numbers are randomly generated and then concatenated using prompt templates for different tasks. The pseudocode is as follows.

Algorithm 1 Generate Arithmetic Dataset (gen_2)

```

309 1: procedure GEN_2(fun, n, d)
310 2:   for digits  $\leftarrow 1$  to d do
311 3:     for i  $\leftarrow 1$  to n do
312 4:       if fun = div then
313 5:         a  $\sim$  Uniform( $10^{2 \cdot \text{digits}-1}$ ,  $10^{2 \cdot \text{digits}} - 1$ )
314 6:         b  $\sim$  Uniform( $10^{\text{digits}-1}$ ,  $10^{\text{digits}} - 1$ )
315 7:       else
316 8:         a  $\sim$  Uniform( $10^{\text{digits}-1}$ ,  $10^{\text{digits}} - 1$ )
317 9:         b  $\sim$  Uniform( $10^{\text{digits}-1}$ ,  $10^{\text{digits}} - 1$ )
318 10:      end if
319 11:      c  $\leftarrow$  fun(a, b)
320 12:      Output sample (digits, a, b, c)
321 13:    end for
322 14:  end for
323 15: end procedure

```

Table 3: Arithmetic Prompts

Task	Prompt
Add	$a + b = ?$
Sub	$a - b = ?$
Mul	$a * b = ?$
Div	Perform integer division: $a/b = ?$
Copy	Copy the following number: a
Rev	Reverse the following number: a
Box	Put the following number in parentheses only: a , example: (number)
Space	Insert a space between every digit in the following number: a
Len	How many digits are in the following number: a , put answer in parentheses only, example: (number)
Count	How many 0 are in the following number: a , put answer in parentheses only, example: (number)
B2d	Convert the following binary number to decimal: a
D2b	Convert the following decimal number to binary: a

4 EXPERIMENTAL RESULTS

4.1 SETUP

We compared several state-of-the-art models:

LLaMA series Dubey et al. (2024), Qwen series Yang et al. (2024; 2025); Team (2025), DeepSeek series Guo et al. (2025), GPT series OpenAI (2023); Hurst et al. (2024)

Both open-source and closed-source models were used to ensure a comprehensive evaluation.

All tasks were tested with problems randomly generated for each digit length from 1 to 100. For Qwen and LLaMA, $n = 10$ problems were generated per digit length. For DeepSeek and GPT, due to resource limitations and slower inference speed, $n = 1$ problem per digit length was used.

Table 4: Model performance on Arithmetic-Bench (Main tasks)

Model	add	sub	mul	div	add_1	sub_1	mul_1	div_1
Llama-3-8B-Instruct	11.7%	10.0%	1.9%	1.8%	93.4%	93.5%	20.3%	18.4%
Llama-3-70B-Instruct	20.5%	16.3%	2.2%	2.2%	93.4%	91.9%	27.2%	26.4%
Qwen2.5-0.5B-Instruct	8.1%	7.3%	1.6%	1.4%	23.1%	19.1%	17.0%	18.1%
Qwen2.5-1.5B-Instruct	12.6%	11.6%	2.0%	1.7%	69.4%	67.8%	27.3%	21.5%
Qwen2.5-3B-Instruct	11.7%	14.1%	1.9%	1.7%	77.6%	71.1%	30.7%	31.5%
Qwen2.5-7B-Instruct	25.7%	20.3%	2.2%	2.9%	89.2%	86.5%	45.2%	45.2%
Qwen2.5-14B-Instruct	28.9%	32.3%	2.3%	3.2%	98.0%	97.7%	66.2%	77.9%
Qwen2.5-32B-Instruct	50.2%	31.3%	2.5%	4.2%	96.9%	97.5%	79.1%	79.0%
Qwen2.5-72B-Instruct	31.5%	30.8%	2.5%	4.3%	98.0%	96.9%	51.5%	48.5%
DeepSeek-R1-Distill-Llama-8B	9.0%	8.0%	3.0%	2.0%	75.0%	65.0%	19.0%	18.0%
DeepSeek-R1-Distill-Llama-70B	14.0%	14.0%	3.0%	5.0%	93.0%	91.0%	27.0%	25.0%
DeepSeek-R1-Distill-Qwen-1.5B	10.0%	12.0%	4.0%	3.0%	44.0%	64.0%	23.0%	21.0%
DeepSeek-R1-Distill-Qwen-7B	14.0%	11.0%	4.0%	3.0%	77.0%	77.0%	32.0%	32.0%
DeepSeek-R1-Distill-Qwen-14B	13.0%	18.0%	4.0%	4.0%	78.0%	77.0%	27.0%	23.0%
DeepSeek-R1-Distill-Qwen-32B	21.0%	26.0%	4.0%	7.0%	83.0%	75.0%	42.0%	43.0%
DeepSeek-R1-671B	46.0%	58.0%	10.0%	10.0%	100.0%	99.0%	56.0%	69.0%
QwQ-32B	26.0%	26.0%	11.0%	10.0%	99.0%	96.0%	41.0%	69.0%
Qwen3-235B-A22B	41.0%	40.0%	10.0%	11.0%	100.0%	100.0%	58.0%	78.0%
gpt-4	51.0%	38.0%	3.0%	4.0%	100.0%	99.0%	61.0%	74.0%
gpt-4o	68.0%	84.0%	3.0%	3.0%	100.0%	100.0%	85.0%	79.0%
gpt-3.5	15.0%	21.0%	3.0%	3.0%	97.0%	89.0%	29.0%	48.0%

378 4.2 ANALYSIS
379380 Main results are shown in Table 4. The accuracy of addition and subtraction is comparable, as is
381 that of multiplication and division. Multiplication is significantly more challenging than addition,
382 with accuracy roughly proportional to the maximum number of digits the model can handle. On
383 multiplication tasks, the best models, Deepseek-R1, QwQ and Qwen3, can correctly solve numbers
384 with up to 10 digits.385 Tasks involving $n \times 1$ -digit numbers are relatively easier, yet most models still fail to achieve 100%
386 accuracy. Tasks where models perform relatively well include **add_1, sub_1, copy, box, and space**,
387 with some models reaching perfect accuracy. These tasks share a common feature: they do not
388 require complex reasoning, and the input-output structures are largely similar. For example, in
389 **add_1** and **sub_1**, changes mostly occur in the last digits.390 **Limitations of Reasoning.** Reasoning models, such as Deepseek-R1, QwQ and Qwen3, outper-
391 form non-reasoning models on tasks like multiplication and base conversion, but underperform on
392 simpler tasks, such as addition and single-digit multiplication. This phenomenon is consistent with
393 the observations reported in Shojaee et al. (2025): on low-complexity tasks, non-reasoning mod-
394 els outperform reasoning models; on medium-complexity tasks, reasoning models demonstrate an
395 advantage; and on high-complexity tasks, both types of models experience complete failure.396 **Influence of Scaling.** Within the same model series, larger models generally perform better on
397 arithmetic tasks. However, Qwen-72B does not outperform Qwen-32B, suggesting that merely in-
398 creasing model size does not necessarily resolve arithmetic challenges.399 **Influence of Distillation.** The DeepSeek distilled models perform worse than their corresponding
400 Qwen counterparts on simple arithmetic tasks, like addition and subtraction, and only marginally
401 outperform Qwen on multiplication and counting. This indicates limitations in their reasoning abil-
402 ity, suggesting that the full reasoning capability of a large model may not have been successfully
403 distilled into these smaller models.404 GPT-4’s average performance falls between Qwen2.5 and Qwen3, indicating that closed-source
405 models do not necessarily demonstrate stronger arithmetic capabilities. Overall, the accuracy of
406 all models remains relatively low, and true generalization in arithmetic has yet to be achieved.408 4.3 LENGTH GENERALIZATION
409410 As shown in Figure 1, accuracy decreases significantly as the number of digits increases. For multi-
411 plication, once the number of digits exceeds a certain threshold, models consistently fail to produce
412 correct results. Therefore, the overall accuracy is approximately equal to the maximum number of
413 digits in multiplication that the model can handle.415 Even the largest and most advanced models, including Qwen3, DeepSeek, and GPT-4, continue to
416 struggle with arithmetic tasks at scale. Specifically, they are unable to reliably solve 10-digit mul-
417 tiplication and often fail at 100-digit addition, despite their strong performance on a wide range
418 of natural language tasks. This indicates that scaling alone does not resolve the fundamental chal-
419 lenges of arithmetic reasoning, and that current architectures still lack robust mechanisms for exact,
420 length-generalizable computation.421 4.4 MEMORIZATION OF FINITE DATASETS
422423 As shown in Figure 2, training on AIME test set can push accuracy to 100%. We also observed sim-
424 ilar phenomena on other finite datasets, demonstrating that finite benchmarks are prone to cheating.
425 Notably, it requires around 100 epochs to memorize well, rather than remembering it after a single
426 pass. Even after reaching 100% accuracy, fluctuations may still occur.428 4.5 CORRELATION BETWEEN REASONING AND ARITHMETIC
429430 As shown in Figure 5, The model’s performance on mathematical benchmarks such as AIME is
431 positively correlated with its performance on large-number multiplication. Reasoning models ex-

hibit stronger multiplication ability compared to non-reasoning models, but perform worse than non-reasoning models on simple tasks such as addition.

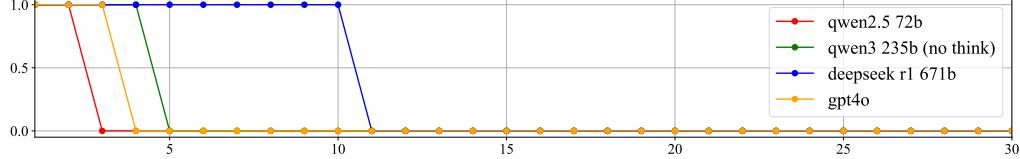


Figure 1: Length Generalization Curve of Multiplication, x axis is length from 1 to 30, y axis is accuracy.

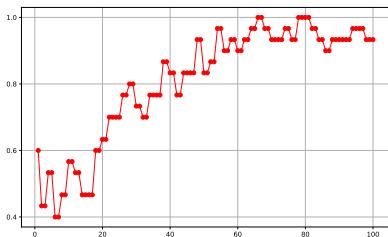


Figure 2: Results of training on AIME 2024, x axis is epoch, y axis is accuracy.

Table 5: Comparison of performance on Multiplication and AIME 2024

Model	Mul Acc (%)	AIME Acc (%)
Qwen2.5 72b	2	13.5
Qwen3 235b (no think)	4	40.1
Qwen3 235b (think)	10	85.7
QwQ	11	79.5
Deepseek r1 671b	10	79.8
gpt4o	3	11.1

4.6 USE OF EXTERNAL TOOL

While it is certainly possible to solve these problems using a calculator Schick et al. (2023)—and, in fact, the web version of ChatGPT often does so, Arithmetic-Bench is fundamentally different. It is designed to use arithmetic as a proxy for abstract reasoning, providing a controlled setting to evaluate a model’s ability to perform multi-step reasoning rather than relying on external tools.

On the other hand, the results of Arithmetic-Bench can be interpreted in two possible ways:

1. In principle, the model’s probabilistic predictions are capable of stable multi-step reasoning, but current models have not realized this ability.
2. The model’s probabilistic predictions cannot guarantee stable multi-step reasoning. If this is the case, it indicates that using external tools for verification is necessary.

4.7 REPRODUCIBILITY

The results from two independently randomly generated sets of problems show little difference, with average fluctuations below 1%. We will make our code publicly available to ensure reproducibility.

5 CONCLUSION

Arithmetic-Bench provides a rigorous, dynamic, and scalable evaluation of LLMs’ multi-step reasoning abilities. Our theoretical analysis shows that the inability to generalize in arithmetic implies broader limitations in general reasoning. Empirical results demonstrate that even state-of-the-art models still struggle with large-number multiplication, highlighting the necessity of improving reasoning mechanisms. We believe that arithmetic-based tasks form the foundation for advancing reasoning in LLMs. Future work should focus on improving data representation, training strategies, verifying, and memory mechanisms. Only by addressing these limitations can LLMs truly perform multi-step reasoning tasks and move beyond shallow inference based on pattern matching.

486 REFERENCES
487

- 488 Zeyuan Allen-Zhu and Yuanzhi Li. Physics of language models: Part 3.3, knowledge capacity
489 scaling laws. *arXiv preprint arXiv:2404.05405*, 2024.
- 490 Yoshua Bengio, Jérôme Louradour, Ronan Collobert, and Jason Weston. Curriculum learning. In
491 *Proceedings of the 26th International Conference on Machine Learning (ICML)*, pp. 41–48, 2009.
492
- 493 Kevin Buzzard and contributors. Ongoing lean formalisation of the proof of fermat’s last theo-
494 rem. <https://github.com/ImperialCollegeLondon/FLT>, 2025. Funded by EP-
495 SRC grant EP/Y022904/1 until September 2029.
- 496 Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser,
497 Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, Christopher Hesse, and John
498 Schulman. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*,
499 2021.
500
- 501 Leonardo de Moura and Sebastian Ullrich. The lean 4 theorem prover and programming language.
502 In *Proceedings of the 28th International Conference on Automated Deduction (CADE-28)*, pp.
503 378–388. Springer, 2021. doi: 10.1007/978-3-030-79876-5_37. URL https://doi.org/10.1007/978-3-030-79876-5_37.
- 504 Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha
505 Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. The llama 3 herd of models.
506 *arXiv e-prints*, pp. arXiv–2407, 2024.
507
- 508 Google. Nano banana (gemini 2.5 flash image) model. Google AI / Gemini Developer API, 2025.
509 URL <https://developers.google.com/>. Image generation / editing model, publicly
510 documented.
511
- 512 Alex Graves, Greg Wayne, and Ivo Danihelka. Neural turing machines. *arXiv preprint*
513 *arXiv:1410.5401*, 2014.
514
- 515 Daya Guo, Dejian Yang, Huawei Zhang, Junxiao Song, Ruoyu Zhang, Runxin Xu, Qihao Zhu,
516 Shirong Ma, Peiyi Wang, Xiao Bi, et al. Deepseek-r1: Incentivizing reasoning capability in llms
517 via reinforcement learning. *arXiv preprint arXiv:2501.12948*, 2025.
- 518 Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob
519 Steinhardt. Measuring mathematical problem solving with the math dataset. In *International
520 Conference on Learning Representations (ICLR)*, 2021.
521
- 522 Aaron Hurst, Adam Lerer, Adam P Goucher, Adam Perelman, Aditya Ramesh, Aidan Clark, AJ Os-
523 trow, Akila Welihinda, Alan Hayes, Alec Radford, et al. Gpt-4o system card. *arXiv preprint*
524 *arXiv:2410.21276*, 2024.
- 525 Eldar Kurtic, Amir Moeini, and Dan Alistarh. Mathador-lm: A dynamic benchmark for mathemati-
526 cal reasoning on large language models. *arXiv preprint arXiv:2406.12572*, 2024.
527
- 528 Nayoung Lee, Kartik Sreenivasan, Jason D Lee, Kangwook Lee, and Dimitris Papailiopoulos.
529 Teaching arithmetic to small transformers. *arXiv preprint arXiv:2307.03381*, 2023.
530
- 531 Tiedong Liu and Bryan Kian Hsiang Low. Goat: Fine-tuned llama outperforms gpt-4 on arithmetic
532 tasks. *arXiv preprint arXiv:2305.14201*, 2023.
- 533 Iman Mirzadeh, Keivan Alizadeh, Hooman Shahrokhi, Oncel Tuzel, Samy Bengio, and Mehrdad
534 Farajtabar. Gsm-symbolic: Understanding the limitations of mathematical reasoning in large
535 language models. *arXiv preprint arXiv:2410.05229*, 2024.
536
- 537 Maxwell I. Nye, Anders Johan Andreassen, Guy Gur-Ari, Henryk Michalewski, Jacob Austin, David
538 Bieber, David Dohan, Aitor Lewkowycz, Maarten Bosma, Charles Sutton, and Augustus Odena.
539 Show your work: Scratchpads for intermediate computation with language models. In *Inter-
540 national Conference on Learning Representations (ICLR)*, 2021.

- 540 OpenAI. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023. URL <https://arxiv.org/abs/2303.08774>.
- 541
- 542
- 543 Long Ouyang, Jeffrey Wu, Xu Jiang, Diogo Almeida, Carroll Wainwright, Pamela Mishkin, Chong
544 Zhang, Sandhini Agarwal, Katarina Slama, Alex Ray, et al. Training language models to fol-
545 low instructions with human feedback. *Advances in neural information processing systems*, 35:
546 27730–27744, 2022.
- 547 Timo Schick, Jane Dwivedi-Yu, Roberto Dessi, Roberta Raileanu, Maria Lomeli, Luke Zettlemoyer,
548 Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to
549 use tools. *arXiv preprint arXiv:2302.04761*, 2023. URL <https://arxiv.org/abs/2302.04761>.
- 550
- 551 Parshin Shojaee, Iman Mirzadeh, Keivan Alizadeh, Maxwell Horton, Samy Bengio, and Mehrdad
552 Farajtabar. The illusion of thinking: Understanding the strengths and limitations of reasoning
553 models via the lens of problem complexity. *arXiv preprint arXiv:2506.06941*, 2025.
- 554
- 555 Aarohi Srivastava, Abhinav Rastogi, Abhishek Rao, et al. Beyond the imitation game: Quantify-
556 ing and extrapolating the capabilities of language models. *Transactions on Machine Learning
557 Research*, 2023. URL <https://openreview.net/forum?id=uyTL5Bvosj>.
- 558 Qwen Team. Qwq-32b: Embracing the power of reinforcement learning, March 2025. URL
559 <https://qwenlm.github.io/blog/qwq-32b/>.
- 560
- 561 Hemish Veeraboina. Aime problem set 1983-2024, 2023. URL <https://www.kaggle.com/datasets/hemishveeraboina/aime-problem-set-1983-2024>.
- 562
- 563 Jason Wei, Xuezhi Wang, Dale Schuurmans, Maarten Bosma, Brian Ichter, Fei Xia, Ed Chi, Quoc
564 Le, and Denny Zhou. Chain-of-thought prompting elicits reasoning in large language models.
565 *arXiv preprint arXiv:2201.11903*, 2022.
- 566
- 567 Tianwen Wei, Jian Luan, Wei Liu, Shuang Dong, and Bin Wang. Cmath: Can your language model
568 pass chinese elementary school math test? *arXiv preprint arXiv:2306.16636*, 2023.
- 569
- 570 An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li,
571 Dayiheng Liu, Fei Huang, Haoran Wei, Huan Lin, Jian Yang, Jianhong Tu, Jianwei Zhang, Jianxin
572 Yang, Jiaxi Yang, Jingren Zhou, Junyang Lin, Kai Dang, Keming Lu, Keqin Bao, Kexin Yang,
573 Le Yu, Mei Li, Mingfeng Xue, Pei Zhang, Qin Zhu, Rui Men, Runji Lin, Tianhao Li, Tianyi Tang,
574 Tingyu Xia, Xingzhang Ren, Xuancheng Ren, Yang Fan, Yang Su, Yichang Zhang, Yu Wan,
575 Yuqiong Liu, Zeyu Cui, Zhenru Zhang, and Zihan Qiu. Qwen2.5 technical report. *arXiv preprint
576 arXiv:2412.15115*, 2024.
- 577
- 578 An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu,
579 Chang Gao, Chengan Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint
580 arXiv:2505.09388*, 2025.
- 581
- 582 Zhen Yang, Ming Ding, Qingsong Lv, Zhihuan Jiang, Zehai He, Yuyi Guo, Jinfeng Bai, and Jie Tang.
583 Gpt can solve mathematical problems without a calculator. *arXiv preprint arXiv:2309.03241*,
584 2023.
- 585
- 586 Zheng Yuan, Hongyi Yuan, Chuanqi Tan, Wei Wang, and Songfang Huang. How well do large
587 language models perform in arithmetic tasks?, 2023.
- 588
- 589 Qihuang Zhong, Kang Wang, Ziyang Xu, Juhua Liu, Liang Ding, and Bo Du. Achieving 97% on
590 gsm8k: Deeply understanding the problems makes llms better solvers for math word problems.
591 *arXiv preprint arXiv:2404.14963*, 2024.
- 592
- 593 Denny Zhou. Llm reasoning (stanford cs-25 lecture). Lecture notes, available at <https://dennyyzhou.github.io/LLM-Reasoning-Stanford-CS-25.pdf>, 2025. Accessed:
2025-09-22.
- 594 Lukasz Kaiser and Ilya Sutskever. Neural gpus learn algorithms. In *International Conference on
595 Learning Representations (ICLR)*, 2016.

594 **A APPENDIX**

595

596 **A.1 IMPLEMENTATION DETAILS**

597

598 We conducted experiments using 1–2 machines equipped with 8×A100 GPUs and deployed our
 599 models based on vLLM. The number of GPUs required varies depending on the model size. To
 600 improve efficiency, we employ parallel inference acceleration for smaller models. Each model was
 601 evaluated using the officially recommended decoding parameters. Ablation studies show that the
 602 decoding parameters have little effect on the results.

603 For reasoning models, context length has a significant impact: if the maximum length is insufficient
 604 to generate the complete output, performance will decrease a lot. Therefore, it is ultimately set to
 605 16,384. For non-reasoning models, since their responses are naturally short. A maximum length of
 606 2,048 or 4,096 is sufficient.

607

608 **A.2 PROOF**

609

610 Proof of Theorem 1

611

612 **Theorem 1.** A container with capacity a cannot hold information larger than a .

613

614 **Proof.** Suppose a container with capacity a_1 can hold information of size a_0 , where $a_1 < a_0$. Then
 615 there exists a container with capacity $a_2 < a_1$ that can hold the a_1 -capacity container. Repeating
 616 this operation, we can construct a decreasing sequence of capacities $a_n < \dots < a_2 < a_1 < a_0$.

617 Since capacities are non-negative, by the monotone bounded sequence theorem, this sequence must
 618 have a limit.

619 **Case 1:** The limit is 0. Then an empty container could hold information of any size, which is
 620 obviously a contradiction.

621 **Case 2:** The limit is greater than 0. Then for each capacity a , there exists a corresponding lower
 622 bound $b < a$ such that a container of capacity b can hold the container of capacity a . Similarly, for
 623 b , there exists a lower bound $c < b$ such that c can hold b , and thus c can hold a . This contradicts
 624 the assumption that b is the lower bound of the sequence of capacities.

625 Therefore, the proposition is proved.

626

627

628 **A.3 SUB TASKS**

629

630 In addition to standard arithmetic operations, we also evaluated several sub-tasks as a complement
 631 to the main tasks. As shown in Table 6. The positive correlation between subtask performance and
 632 the main task indicates that proficiency on subtasks reflects or contributes to overall performance on
 633 the main task.

634 Most models struggle with sub-tasks such as reversing and counting. Small models (0.5B and 1.5B)
 635 exhibit clear flaws in instruction-following, showing low accuracy on tasks like **copy**, **box**, and
 636 **space**.

637

638 **A.4 EXPERIMENTS OF MEMORIZATION**

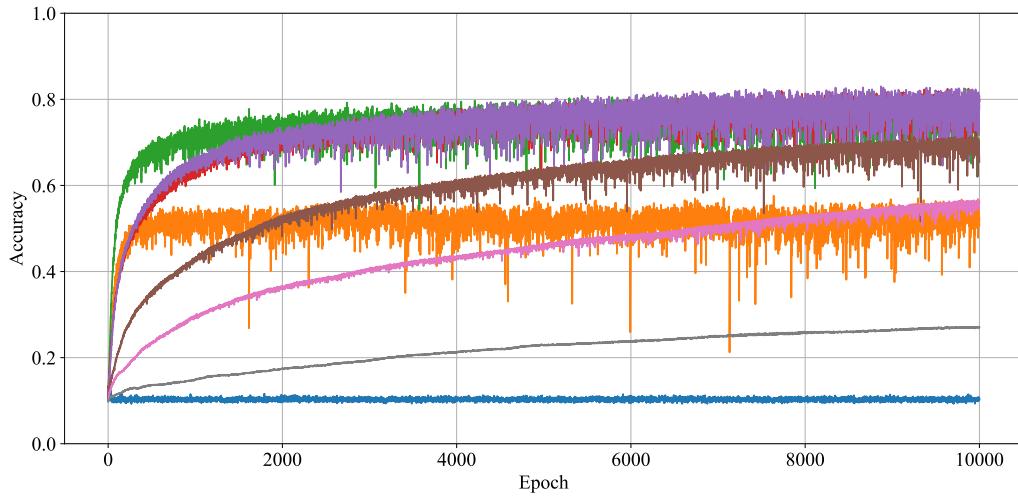
639

640 We constructed a dataset using the first 10,000 digits of π , where the input is the index (the n -th
 641 digit) and the output is the corresponding digit represented as a one-hot vector. Models with varying
 642 parameter sizes were trained to memorize the π data up to their capacity limit, defined as the point
 643 where the accuracy converges. The learning rate was optimized via grid search to maximize the
 644 converged accuracy. We conducted dozens of experiments, and the result of one representative run
 645 is shown in Figure 3. The curves from the other experiments exhibit similar shapes.

646 For different model, we calculated the ratio of the information content of the correctly memorized
 647 digits to the total number of model parameters. Our experiments show that this ratio is quite stable,
 648 regardless of model size or the number of digits, and is approximately 2.2 bits per parameter.

Table 6: Model performance on Arithmetic-Bench (Sub tasks)

Model	copy	rev	box	space	length	count	b2d	d2b
Llama-3-8B-Instruct	100.0%	3.7%	99.8%	30.2%	8.6%	14.4%	5.5%	1.9%
Llama-3-70B-Instruct	100.0%	7.1%	99.9%	80.7%	11.7%	6.5%	10.1%	3.1%
Qwen2.5-0.5B-Instruct	69.2%	2.8%	93.1%	20.3%	0.2%	7.8%	5.0%	1.3%
Qwen2.5-1.5B-Instruct	69.1%	9.3%	99.9%	82.5%	6.7%	6.2%	2.1%	3.1%
Qwen2.5-3B-Instruct	99.9%	13.4%	99.5%	86.5%	9.6%	4.3%	8.1%	2.2%
Qwen2.5-7B-Instruct	99.9%	13.4%	99.9%	99.9%	13.0%	24.7%	12.7%	3.5%
Qwen2.5-14B-Instruct	99.9%	17.1%	100.0%	100.0%	15.1%	24.7%	12.7%	3.5%
Qwen2.5-32B-Instruct	100.0%	23.9%	100.0%	99.9%	16.6%	42.9%	15.4%	4.6%
Qwen2.5-72B-Instruct	99.9%	14.2%	99.7%	100.0%	23.3%	32.6%	15.6%	5.0%
DeepSeek-R1-Distill-Llama-8B	99.0%	6.6%	94.0%	37.0%	31.0%	47.0%	9.0%	1.0%
DeepSeek-R1-Distill-Llama-70B	100.0%	9.0%	100.0%	74.0%	33.0%	35.0%	10.0%	3.0%
DeepSeek-R1-Distill-Qwen-1.5B	96.0%	10.0%	77.0%	17.0%	34.0%	13.0%	3.0%	2.0%
DeepSeek-R1-Distill-Qwen-7B	100.0%	14.0%	91.0%	93.0%	36.0%	44.0%	20.0%	5.0%
DeepSeek-R1-Distill-Qwen-14B	100.0%	24.0%	100.0%	100.0%	38.0%	59.0%	11.0%	2.0%
DeepSeek-R1-Distill-Qwen-32B	100.0%	23.0%	100.0%	100.0%	25.0%	42.0%	13.0%	4.0%
QwQ-32B	100.0%	70.0%	99.0%	100.0%	98.0%	99.0%	31.0%	14.0%
Qwen3-235B-A22B	100.0%	78.0%	100.0%	100.0%	100.0%	59.0%	20.2%	
deepseek r1 671b	100.0%	82.0%	100.0%	100.0%	96.0%	100.0%	55.0%	16.0%
gpt4	100.0%	15.0%	100.0%	100.0%	54.0%	22.0%	11.0%	3.0%
gpt4o	100.0%	27.0%	100.0%	100.0%	68.0%	11.0%	11.0%	4.0%
gpt3.5	100.0%	20.0%	90.0%	51.0%	17.0%	10.0%	3.0%	2.0%

Figure 3: Memorization of π , different colors represent different learning rates.

702 A.5 EXPERIMENTS OF FORGETTING
703704 We constructed datasets using the first 10,000 digits of π and e . The model was first trained to
705 memorize the π dataset up to its capacity limit, and then trained on the e dataset. In the input
706 vectors, the indices of π digits were placed on the left, while the indices of e digits were placed on
707 the right, ensuring that the inputs did not overlap.708 Despite the absence of input conflicts, the model completely forgot the π data after learning e . This
709 demonstrates that once a model reaches its capacity limit, adding new information inevitably causes
710 it to forget previously memorized information.
711712 A.6 USE OF AI ASSISTANTS
713714 To reduce the cost of manual revisions, we used ChatGPT Ouyang et al. (2022) to revise the language
715 of the paper. The revisions were made solely to enhance the clarity and readability of the text and
716 not for any other purpose.
717
718
719
720
721
722
723
724
725
726
727
728
729
730
731
732
733
734
735
736
737
738
739
740
741
742
743
744
745
746
747
748
749
750
751
752
753
754
755